

# Real-Time Big Data Processing

## Lab 2 – Processing Real-Time Data with Stream Analytics

### Overview

In this lab, you will create an Azure Stream Analytics job to process the simulated device data you generated in lab 1.

### What You'll Need

To complete the labs, you will need the following:

- A web browser
- A Microsoft account
- A Microsoft Azure subscription
- A Windows, Linux, or Mac OS X computer
- The lab files for this course
- The Azure resources created in the previous labs

**Important:** If you have not completed lab 1, or you have deleted the event hub, storage account, and IoT hub you created, complete lab 1 now.

### Processing a Stream of Data

To process real-time data as it arrives in an event hub or IoT hub, you can use an Azure Stream Analytics job. In this procedure, you will create a simple Stream Analytics job that reads device readings from your event hub, and stores them in a blob store container.

### Create a Stream Analytic Job

The first step in using Stream Analytics to process real-time data is to create a Stream Analytics job.

1. In a web browser, navigate to <http://portal.azure.com>, and if prompted, sign in using the Microsoft account that is associated with your Azure subscription.
2. In the Microsoft Azure portal, in the Hub Menu, click **New**. Then in the **Internet of Things** menu, click **Stream Analytics job**.
3. In the **New Stream Analytics Job** blade, enter the following settings, and then click **Create**:
  - **Name:** Enter a unique name (and make a note of it!)
  - **Subscription:** Select your Azure subscription
  - **Resource Group:** Select the resource group containing your existing resources
  - **Location:** Select any available region
  - **Pin to dashboard:** Not selected

4. In the Azure portal, view **Notifications** to verify that deployment has started. Then wait for the job to be deployed (this can take a few minutes.)

### Add an Input

Stream Analytics jobs get their data from one or more *inputs*. In this procedure, you will create and sample an input for the event hub you created previously.

1. In the Azure portal, browse to the Stream Analytics job you created previously.
2. In the blade for your Stream Analytics job, in the **Job Topology** section, click the **Inputs** tile.
3. In the **Inputs** blade, click **Add**.
5. In the **New input** blade, enter the following settings, and then click **Create**:
  - **Input alias:** DeviceData
  - **Source Type:** Data stream
  - **Source:** IoT hub
  - **Subscription:** Use IoT hub from current subscription
  - **IoT hub:** *Select your IoT hub*
  - **Endpoint:** Messaging
  - **Shared Access Policy Name:** iothubowner
  - **Consumer group:** \$Default
  - **Event serialization format:** JSON
  - **Encoding:** UTF-8
4. Wait for the input to be created and tested.

### Add an Output

Stream Analytics jobs return their results to an *output*. In this procedure, you will add an output to your Stream Analytics job so that the processed results are stored in Azure Blob storage.

1. In the Azure portal, browse to the Stream Analytics job you created previously.
2. In the blade for your Stream Analytics job, in the **Job Topology** section, click the **Outputs** tile.
3. In the **Outputs** blade, click **Add**.
4. In the **New output** blade, enter the following settings, and then click **Create**:
  - **Output alias:** DeviceReadings
  - **Sink:** Blob storage
  - **Subscription:** Use blob storage from current subscription
  - **Storage account:** *Select your storage account*
  - **Container:** Create a new container
  - **Container:** device-readings
  - **Path pattern:** readings/{date}
  - **Date format:** YYYY/MM/DD
  - **Time format:** *Should be unavailable*
  - **Event serialization format:** CSV
  - **Delimiter:** comma (,)
  - **Encoding:** UTF-8
5. Wait for the output to be created and tested.

### Add a Query

Now that you have defined an input and an output for your Stream Analytics job, you can connect them by defining a query that will process the data stream.

1. In the Azure portal, browse to the Stream Analytics job you created previously.

2. In the blade for your Stream Analytics job, in the **Job Topology** section, click the **Query** tile.
3. In the query blade, modify the default query that is provided for you, replacing **YourOutputAlias** with the output alias you specified for your output, and **YourInputAlias** with the input alias you specified for your input (the available inputs and outputs are shown on the left of the query editor pane):

```
SELECT
    *
INTO
    [ DeviceReadings ]
FROM
    [ DeviceData ]
```

4. Save the query, and then close the query pane.

### View the Job Diagram

Your Stream Analytics job now consists of an input, connected to an output by a query. You can verify this by viewing the job diagram.

1. In the blade for your Stream Analytics job, click **Settings**.
2. In the **Settings** blade, click **Job diagram**.
3. In the **Job diagram** blade, verify that your job consists of an IoT Hub input, followed by a query step, followed by a Blob Storage output.
4. Close the **Job diagram** blade and the **Settings** blade.

### Start the Job

A Stream Analytics job runs perpetually, processing data as it arrives.

1. In the blade for your Stream Analytics job, click **Start**. Then in the **Start job** blade, ensure that **Now** is selected and click **Start**.
2. Wait for the streaming job to start – this can take a minute or so.
3. When the job has started, in the Node.JS console, in the **iotdevice** folder, enter the following command to run device simulation script and start submitting messages to the IoT hub:

```
node iotdevice.js
```

4. While the script is running, start Azure Storage Explorer, and if necessary, sign into your azure subscription using your Microsoft account.
5. Expand your storage account, and then expand **Blob Containers**.
6. Double-click the **device-readings** container, and then browse through the **readings** folder, and the year, month, and date folder to view the most recent blob that has been generated by your job.
7. Download the blob to open it in a text editor or spreadsheet application, and verify that in addition to **device**, **reading** fields, it contains values for **EventProcessedUtcTime**, **PartitionId**, **EventEnqueuedUtcTime**, and **IoT Hub**.
8. Close the downloaded file, and in the Node.JS console, press CTRL+C to stop the script.

### Stop the Job

When you want to stop processing events, you can stop the job.

1. In the blade for your Stream Analytics job, click **Stop**. When prompted to confirm, click **Yes**.
2. Wait for the job to stop running.

## Extending a Streaming Solution

A Stream Analytics job can include multiple inputs and outputs, enabling you to combine data streams from multiple sources and send processed results to multiple destinations. Additionally, you can create a streaming topology that uses multiple Stream Analytics jobs to filter and route messages through event hubs. In this exercise, you will extend your streaming solution to filter readings greater than 0.5 and route them to a second Stream Analytics job, which will store them in blob storage in an alerts folder.

### Add a Second Output

In this procedure, you will add a second output to your Stream Analytics job.

1. In the Azure portal, browse to the Stream Analytics job you created previously.
2. In the blade for your Stream Analytics job, in the **Job Topology** section, click the **Outputs** tile.
3. In the **Outputs** blade, click **Add**.
4. In the **New output** blade, enter the following settings, and then click **Create**:
  - **Output alias:** HighReadings
  - **Sink:** Event Hub
  - **Subscription:** Use event hub from current subscription
  - **Service bus namespace:** *Select your service bus namespace*
  - **Event hub name:** *Select your event hub*
  - **Event hub policy name:** *DeviceAccess (this should be the name of the shared access policy you created in lab 1)*
  - **Partition key column:** *Leave blank*
  - **Event serialization format:** JSON
  - **Encoding:** UTF-8
  - **Format:** Line separated
5. Wait for the output to be created and tested.

### Modify the Query

Now that you have defined a second output for your Stream Analytics job, you can send processed data to it from your query.

1. In the Azure portal, browse to the Stream Analytics job you created previously.
2. In the blade for your Stream Analytics job, in the **Job Topology** section, click the **Query** tile.
3. Modify the query as shown below, creating a common table expression for all readings, a SELECT statement that routes the **device**, **reading**, and **EventEnqueuedUtcTime** fields to the blob store output, and a second SELECT statement that filters rows with a reading greater than 0.5 and writes them to the new event hub output:

```
WITH [AllReadings] AS
(SELECT * FROM [DeviceData])

SELECT device, reading, EventEnqueuedUtcTime
INTO [DeviceReadings]
FROM [AllReadings]

SELECT device, reading, EventEnqueuedUtcTime
INTO [HighReadings]
FROM [AllReadings]
WHERE CAST(reading AS float) > 0.5
```

4. Save the query, and then close the query pane.

## View the Job Diagram

Your Stream Analytics job now consists of an input, connected to two outputs by two steps. You can verify this by viewing the job diagram.

1. In the blade for your Stream Analytics job, click **Settings**.
2. In the **Settings** blade, click **Job diagram**.
3. In the **Job diagram** blade, verify that your job consists of an IoT Hub input, followed by two query steps, followed by a Blob Storage output and an Event Hub output.
4. Close the **Job diagram** blade and the **Settings** blade.

## Add a Second Stream Analytics Job

The Stream Analytics job you have created routes readings with a high value to an event hub. You will now add a second stream analytics job to process these high readings.

1. In the Microsoft Azure portal, in the Hub Menu, click **New**. Then in the **Internet of Things** menu, click **Stream Analytics job**.
2. In the **New Stream Analytics Job** blade, enter the following settings, and then click **Create**:
  - **Name**: *Enter a unique name (and make a note of it!)*
  - **Subscription**: *Select your Azure subscription*
  - **Resource Group**: *Select the resource group containing your existing resources*
  - **Location**: *Select any available region*
  - **Pin to dashboard**: *Not selected*
3. In the Azure portal, view **Notifications** to verify that deployment has started. Then wait for the job to be deployed (this can take a few minutes.) Then browse to the blade for the new Stream Analytics job.
4. Add an input to the new job, with the following settings:
  - **Input alias**: HighReadings
  - **Source Type**: Data stream
  - **Source**: Event hub
  - **Subscription**: Use event hub from current subscription
  - **Service bus namespace**: *Select your service bus namespace*
  - **Event hub name**: *Select your event hub*
  - **Event hub policy name**: *DeviceAccess (this should be the name of the shared access policy you created in lab 1)*
  - **Event hub consumer group**: *Leave blank*
  - **Event serialization format**: JSON
  - **Encoding**: UTF-8
6. Add an output to the stream analytics job, with the following settings:
  - **Output alias**: DeviceAlerts
  - **Sink**: Blob storage
  - **Subscription**: Use blob storage from current subscription
  - **Storage account**: *Select your storage account*
  - **Container**: device-readings
  - **Path pattern**: alerts/{date}
  - **Date format**: YYYY/MM/DD
  - **Time format**: *Should be unavailable*
  - **Event serialization format**: CSV
  - **Delimiter**: comma (,)
  - **Encoding**: UTF-8
5. Add the following query to the stream analytics job:

```
SELECT
    *
INTO
    [ DeviceAlerts ]
FROM
    [ HighReadings ]
```

## Start the Jobs

Now you're ready to start both jobs and test the streaming topology.

1. Start the new stream analytics job and wait for it to start – this can take a minute or so.
2. Start the original stream analytics job and wait for it to start – this can take a minute or so.
3. When the jobs have started, in the Node.JS console, in the **iotdevice** folder, enter the following command to run device simulation script and start submitting messages to the IoT hub:

```
node iotdevice.js
```

4. While the script is running, start Azure Storage Explorer, and if necessary, sign into your azure subscription using your Microsoft account.
5. Expand your storage account, and then expand **Blob Containers**.
6. Double-click the **device-readings** container, and then browse through the **alerts** folder, and the year, month, and date folder to view the most recent blob that has been generated by your job.
7. Download the blob to open it in a text editor or spreadsheet application, and verify that it contains readings with a value greater than 0.5.
8. Close the downloaded file, and in the Node.JS console, press CTRL+C to stop the script.

## Stop the Jobs

When you want to stop processing events, you can stop the jobs.

1. In the Azure portal, stop both stream analytics jobs.

**Note:** You will use the resources you created in this lab when performing the next lab, so do not delete them. Ensure that all stream analytics jobs and Node.js scripts are stopped to minimize ongoing resource usage costs.